Advanced AI Skills

Verification, Governance and the Real World

Real Minds AI

Tracy Anthony and Dr Dennis Wollersheim

Why "Advanced"?

The morning exercises used clean data. Real life is messy.

In the real world, the data you feed AI is:

- **Incomplete** – fields missing, logs partial, records half-finished
- **Contradictory** – two people remember the same event differently
- **Ambiguous** – "soon", "when convenient", "where appropriate"
- **Risky** – one wrong word in an email can breach privacy or over-promise

This session teaches you to use AI safely when the inputs are imperfect

> *The question is not "can AI write this?" but "should I send what AI wrote?"*

The AI Sandwich

The workflow that keeps you safe

Step 1: System A outputs (raw notes, exports, logs, emails)
These are messy, partial, sometimes contradictory

Step 2: AI transforms (triage, structure, draft comms)
AI imposes structure and generates drafts – fast but unchecked

Step 3: Human verifies (facts, privacy, policy, approvals)
You check every claim, every name, every commitment

Step 4: System B inputs (incident record, complaint register, final comms)
Clean, verified, auditable output goes into your systems

> *AI sits in the middle.* ***You are the bread.*** *Nothing gets through without your verification.*

Advanced Exercise Overview

#	Exercise	What you will do	Key skill
A1	**Incident Pack Assembly**	Reconcile 2 conflicting accounts into one record	Verification
A2	**Fix the Problematic Email**	Rewrite a privacy-breaching draft	Governance
A3	**Attendance Triage Under Uncertainty**	Build contact plan from messy data	Uncertainty
A4	**Policy Alignment Change Log**	Compare old procedure to new requirements	Compliance
A5	**Messy Minutes to Action Register**	Extract clean actions from vague notes	Structure

Your data pack

The **Messy Mini-Pack** is intentionally broken: duplicate rows, conflicting times, missing fields, privacy traps, and unclear ownership. This is realistic.

Exercise A1

Incident Pack Assembly

Verification and Cross-Checking

The theory: when sources disagree

- In any incident, different people observe different things at different times
- Two staff accounts of the same event **will** contain discrepancies (times, sequence, severity)
- Your job is not to pick one account – it is to **reconcile** them into a single factual record

The Known / Unconfirmed / Unknown framework

Category	Definition	Example
Known	Confirmed by 2+ sources or physical evidence	"S1188 had a graze – confirmed by first aid log"
Unconfirmed	Mentioned by one source, not contradicted	"Bystander called out a name – one account only"
Unknown	Key fact that no source provides	"Exact start time of the altercation"

> *AI can build this framework from messy inputs. **You** decide what moves from "Unconfirmed" to "Known".*

Exercise A1: Incident Pack Assembly

Instructions (20 minutes)

You have **5 documents** about the same incident. Open from your Messy Mini-Pack:

- Artifact 02: Yard duty staff account
- Artifact 03: Student services intake note
- Artifact 04: First aid log
- Artifact 05: Parent contact log
- Artifact 06: CCTV maintenance ticket

Paste your context card, then **all five documents**, then add:

> *Using these five documents about the same incident, please:*
> *1. Create a Known / Unconfirmed / Unknown fact list (bullet points)*
> *2. Produce a SafeTrack incident entry in structured fields: Date, TimeRange, Location, Category, PeopleInvolved (StudentIDs only), ImmediateActions, Injuries, ParentContactAttempts, Evidence, FollowUpOwner, ReviewDate*
> *3. List exactly what must be verified by a human before finalising*

Exercise A1: What You Learned

Key contradictions in the data

- **Time:** Yard duty says "approx 13:05" arrival, Student Services says students arrived "approx 13:12" – times are estimates
- **Physical contact:** Yard duty is uncertain about a punch; both students told Student Services there was no punching
- **First aid log:** Two entries for S1188 at different times (13:16 and 13:18) – possible duplicate or two visits
- **CCTV:** Dashboard showed "NO SIGNAL" during the exact window (13:10-13:25) – footage may not exist
- **Parent contact:** S1048 contact logged at 13:25 (attempt) and 13:32 (success) – Student Services notes "13:25" but may be the attempt time

The principle

> *AI organises conflicting information. **You** verify which facts are confirmed. Never let AI resolve contradictions – that is a human judgment call.*

Exercise A2

Fix the Problematic Email

Spot the Problems

This draft email contains at least 6 governance failures

Read the **Problematic Draft Email** (Artifact 07 in your Messy Mini-Pack):

Problem	Why it matters
Mentions "the other student" and implies identity	**Privacy breach** – must not disclose info about other students
Says student "has a history of behaviour issues"	**Privacy breach** – disclosing another student's record
Claims "we will be suspending them for 5 days"	**Over-promise** – outcome not yet decided or approved
Claims "CCTV clearly shows this"	**False claim** – CCTV had "no signal" during the incident
Promises "this will never happen again"	**Over-promise** – cannot guarantee future behaviour
Asks parent for "screenshots of the group chat"	**Evidence handling risk** – may breach other students' privacy

> *Every one of these could lead to a formal complaint, a privacy investigation, or a legal problem.*

Exercise A2: Fix the Problematic Email

Instructions (15 minutes)

Paste your context card, then the **Problematic Draft Email** (Artifact 07), then add:

> *This draft email is problematic and must not be sent as-is. Rewrite it so that it:*
> *- does not mention or identify any other student*
> *- does not promise specific outcomes (no "never again", no suspension details)*
> *- does not claim evidence that has not been verified*
> *- includes clear next steps with a timeframe and a named contact point*
> *- maintains a calm, empathetic, factual tone*
>
> *Provide three versions:*
> *A. A phone call script (dot points for the call)*
> *B. A follow-up email (under 200 words)*
> *C. An SMS (under 160 characters)*

What to check

- Has the AI removed **all** references to other students?
- Are there any promises of specific outcomes?

Exercise A2: What You Learned

The comms safety checklist

Before sending any incident-related communication, check:

- **No other students identified** (directly or by implication)
- **No outcomes promised** until decisions are formally made
- **No evidence claimed** unless verified and available
- **No blame assigned** in either direction
- **Next steps are specific** (who to contact, by when, what happens next)
- **Tone is calm and factual** (not defensive, not dismissive)

The principle

> *AI drafts fast. Fast drafts are dangerous. The **human review step** is not optional – it is where governance happens.*

Try this at work

- Run every sensitive email through AI with the instruction "check this for privacy breaches, over-promises, and unverified claims"

Exercise A3

Attendance Triage Under Uncertainty

Working with Imperfect Data

The theory: real exports are never clean

The attendance CSV in the Messy Mini-Pack has deliberate problems:

- **Duplicate rows** – S1188 appears twice for the same date
- **Inconsistent codes** – "Unexplained", "LATE", "X" instead of standard U, L codes
- **Missing fields** – blank YearLevel, blank AttendanceCode, blank Prior20Days count
- **Impossible values** – MinutesLate of -3 (negative time)
- **Typos** – "S12O2" instead of "S1202" (letter O instead of zero)
- **Date format inconsistency** – one row uses DD/MM/YYYY instead of YYYY-MM-DD
- **Non-standard homeroom** – "8Z" does not exist in the school

What this exercise demonstrates

- AI can **identify data quality issues** before you start analysis
- Cleaning data is a necessary step before making decisions
- AI handles the tedious work of spotting inconsistencies – you make the calls

Exercise A3: Attendance Triage

Instructions (20 minutes)

Paste the **Messy Attendance CSV** (Artifact 01 from Messy Mini-Pack), then add:

> *This attendance export has data quality issues. Please:*
> *1. List every data quality problem you find (duplicates, inconsistent codes, missing fields, impossible values, typos, format issues)*
> *2. For each problem, suggest how to fix it or what to verify*
> *3. After noting the issues, produce a triage list: students who need parent contact this week, ranked by urgency (highest absence/late pattern first)*
> *4. For each student on the triage list, draft a one-sentence contact note appropriate to their situation*
>
> *Where data is uncertain, note what you are assuming and what needs to be verified in the source system.*

What to look for

- Does AI catch the **S12O2 typo** (letter O instead of zero)?
- Does it flag the **negative minutes late** value?
- Does it identify the **duplicate S1188 row**?

Exercise A3: What You Learned

Data quality issues in the messy export

Issue	Row	Detail
Duplicate	S1188, 2026-03-04	Exact duplicate row
Code inconsistency	S1019, 2026-03-04	"Unexplained" instead of "U"
Code inconsistency	S1007, 2026-03-04	"LATE" instead of "L"
Unknown code	S1089, 2026-03-05	"X" is not a valid attendance code
Impossible value	S1102, 2026-03-04	MinutesLate = -3
Missing field	S1095, 2026-03-04	YearLevel is blank
Typo (StudentID)	2026-03-05	"S12O2" should be "S1202"
Date format	S1048, 2026-03-05	"05/03/2026" instead of YYYY-MM-DD
Invalid homeroom	S1048, 2026-03-06	"8Z" does not exist
Missing code	S1121, 2026-03-05	AttendanceCode is blank

The principle

*Clean data before you analyse. AI finds the mess. **You** decide the fixes.*

Exercise A4

Policy Alignment Change Log

AI as Compliance Translator

The theory: old policy vs new requirements

- Schools receive new regulatory bulletins regularly
- Comparing a 3-page local procedure against a 2-page authority bulletin **by hand** takes hours
- AI can do a structured comparison in seconds – flagging gaps, risks, and required updates

The "change log" pattern

Ask AI to produce a table with:

Column	Purpose
Requirement	What the new bulletin requires
Current procedure	What your local document says
Gap or risk	Where they do not align
Required update	What needs to change
Owner	Who should make the change
Evidence	How you prove compliance

This pattern works for **any** regulatory comparison – not just incident recording.

Exercise A4: Policy Alignment

Instructions (15 minutes)

Paste the **Old Local Procedure** (Artifact 09) AND the **Authority Bulletin** (Artifact 10), then add:

> *Compare the school's local procedure with the authority bulletin.*
> *Output a change log table with columns: Requirement, CurrentProcedure, Gap or Risk, RequiredUpdate, OwnerRole, Evidence.*
> *Then draft a revised 10-14 step procedure that meets the new minimum expectations. Include owners and timeframes for each step.*

Key gaps to expect

- **Recording timeframe:** Local says "within a few days" vs authority says "within 24 hours"
- **Complaints acknowledgement:** Local says "5 business days" vs good practice is 2
- **Unconfirmed marking:** Authority requires "UNCONFIRMED" labels; local procedure has no such requirement
- **Privacy guidance:** Local says "provide enough detail"; authority says "do not disclose other students"

What to check

- Has AI identified all the major gaps?

Exercise A4: What You Learned

Key takeaways

- AI performs **structured document comparison** faster and more thoroughly than manual review
- The change log format makes gaps visible and actionable – not buried in paragraph text
- AI can draft a revised procedure, but the **owner assignments and timeframes** need human input

The principle

> *AI translates between "what the authority requires" and "what your school currently does" – the gap analysis is instant, the implementation is human*

Try this at work

- Every time a new circular, bulletin, or regulatory update arrives, paste it alongside your current procedure
- Ask for a change log table – this becomes your compliance evidence
- Assign owners in a staff meeting, not in the AI output

Exercise A5

Messy Minutes to Action Register

Imposing Structure on Chaos

The theory: meeting notes are always messy

The messy exec meeting notes (Artifact 11) have deliberate problems:

- **Unclear ownership** – "send list to coordinators" (who sends? which coordinators?)
- **Vague dates** – "maybe Wed?", "due Tue (or Wed?)", "start soon"
- **Missing decisions** – "decision-ish" (was it actually decided?)
- **Ambiguous actions** – "need a plan" (what plan? who writes it?)

What this exercise demonstrates

- AI can extract action items even from terrible notes
- AI **flags ambiguity** rather than silently resolving it
- The output is a clean action register that you can paste into your task system
- The human step is confirming owners, dates, and whether "decision-ish" items were actually decided

Exercise A5: Messy Minutes

Instructions (15 minutes)

Paste the **Messy Exec Meeting Notes** (Artifact 11), then add:

> *These meeting notes are messy and incomplete. Please:*
> *1. Format as professional meeting minutes with clear headings*
> *2. Extract all action items into a table: Action, Owner, DueDate, Status*
> *3. Where the owner is unclear, write "OWNER TBC" and suggest who it might be*
> *4. Where the date is vague, write "DATE TBC" and suggest a reasonable deadline*
> *5. List any items that sound like decisions but are not clearly confirmed – mark these as "DECISION TBC"*
> *6. Draft a 5-sentence email summary to send to attendees, asking them to confirm their action items*

What to look for

- Does AI flag "maybe Wed?" as uncertain?
- Does it handle "decision-ish" appropriately?
- Does the email ask attendees to **confirm**, not just inform?

Exercise A5: What You Learned

Key ambiguities in the messy notes

Item	Problem	Resolution needed
"Send list to coordinators"	Who sends? Which coordinators?	Owner and recipient TBC
"Due Tue (or Wed?)"	Conflicting dates	Confirm with Wellbeing Leader
"Decision-ish: increase supervision"	Was this formally decided?	Confirm with Principal
"Trial weekly digest – start soon"	No start date, no owner	Assign owner and date
"10-slide deck outline – who drafts?"	No owner assigned	Assign at next meeting

The principle

> *AI imposes **structure** on messy inputs. You impose **accountability** – confirming owners, dates, and decisions. Neither step works without the other.*

The Advanced Toolkit

Five patterns for messy data

Pattern	When to use	Key output
Known / Unconfirmed / Unknown	Conflicting sources	Fact classification list
Comms Safety Check	Any sensitive draft	Privacy and promise audit
Data Quality Audit	Messy exports / CSVs	Issue list with fix suggestions
Change Log Comparison	Old vs new requirements	Gap table with owners
Ambiguity Flagging	Vague meeting notes	Action register with TBCs

The golden rule for all advanced exercises

> *AI is fast. Fast is dangerous when the data is messy.* ***The human verification step is where safety happens.***

Bringing It Together

The AI Sandwich in practice

Every task you do at work follows this pattern:

Raw input (messy email, partial export, scribbled notes)

↓

AI transforms (structures, drafts, flags issues)

↓

You verify (check facts, check privacy, check promises)

↓

Clean output (system entry, sent email, filed record)

Your verification checklist

- Are **facts** confirmed by 2+ sources?
- Are **people** identified only by role or ID (not by name in general text)?
- Are **promises** limited to what has been decided?
- Are **evidence** claims backed by actual records?
- Are **owners and dates** specific (not "someone" or "soon")?

Thank You

Tracy Anthony and **Dr Dennis Wollersheim**
Real Minds AI

hello@realmindsai.com.au | realmindsai.com.au

wisdom, amplified